



Exercise Sheet 7

1. Exercise: Linear Classifier for Multiple Classes (MNIST)

- (a) Load the digits dataset provided by the scikit-learn (`sklearn`) library and generate the following subsets. Note that each subset consists of data samples and their corresponding labels.
 - (i) `train_data`: The training subset of the dataset makes up 70% of the original dataset. This is the largest portion of the dataset and it's used directly for training the model.
 - (ii) `test_data`: The test subset of the dataset makes up 30% of the original dataset. The test dataset is a subset used to provide an unbiased evaluation of a final trained model. It is used to verify the accuracy of the model built.
- (b) Implement the t-SNE algorithm to reduce the data dimensionality down to four dimensions ($f_{\text{t-SNE}} : \mathbb{R}^{64} \rightarrow \mathbb{R}^4$). Apply this procedure for the training, test, and validation datasets.
- (c) Implement the *One-versus-the-rest* multiclass strategy using K linear dichotomy-classifiers. Write a Python class, called `OneVsRest`, that implements the *One-versus-the-rest* multiclass classifier and should include the following methods:
 - (i) `__init__`: This function initializes the classifier given the number K of dichotomy-classifiers.
 - (ii) `fit`: This should fit the model to the provided training data.
 - (iii) `predict`: This should predict class labels for the provided test data samples.
- (d) Utilize the *One-versus-the-rest* classifier in order to classify the data samples of the digit test dataset. Therefore, initialize the class using $K = 10$ and fit the classifier using the dimensionality reduced train data (\mathbb{R}^4). Utilize the trained model to make predictions on the labels of the (dimensionality reduced) test dataset.
- (e) Evaluate the performance of the implemented classifier. Therefore, write a Python function called `evaluate(y_pred, y_gt)` to determine the accuracy of a classification model's predictions. The function computes the accuracy as the proportion of correct predictions over the total number of predictions.
- (f) Apply the t-SNE algorithm to reduce the data dimensionality of the original test dataset down to two dimensions ($f_{\text{t-SNE}} : \mathbb{R}^{64} \rightarrow \mathbb{R}^2$). Generate one figure with two distinct subplots that show the dimensionally reduced test dataset, using the predicted labels from the *One-versus-the-rest* classifier and the actual ground truth labels, respectively.

2. Exercise: Softmax Classification

This exercise focuses on the implementation of a softmax classifier using the PyTorch library. It provides a way to understand the functionalities and potential applications of PyTorch, a popular library in the field of machine learning. To implement and use the classifier, follow these steps:

(a) **Data Preprocessing:**

Use the dataset created in Exercise 1 (a). Transform the target classes of the train and test datasets into one-hot encoded class labels.

(b) **Model Definition and Initialization:**

Using PyTorch, define a model by creating a class that inherits from the Module class. To do so, start by importing the PyTorch library using:

```
import torch as torch
```

and generate the class `SoftmaxClassifier(torch.nn.Module)` with the functions:

- (i) `__init__(self, n_feat, n_classes)`: This function defines the layers of the classification network. The model should be composed of one linear layer and one softmax function. To define the layers use:

```
super().__init__()
self.linear = torch.nn.Linear(n_inputs, n_outputs)
self.softmax = softmax() %TODO
```

Implement the `softmax()` function without using built-in functionality from a library, computing

$$P(c_l|\mathbf{x}) = \frac{e^{a_l}}{\sum_{k=1}^K e^{a_k}}, \quad (1)$$

where $\mathbf{a} = \boldsymbol{\theta}_k^T \tilde{\mathbf{x}}$ is the output of the linear transformation of data $\tilde{\mathbf{x}} = [\mathbf{x}^T \ 1]^T$. Note, that the linear layer `torch.nn.Linear()` intrinsically performs the mapping $\mathbf{x} \rightarrow \tilde{\mathbf{x}}$.

- (ii) `forward(self, x)`: Defines how the data \mathbf{x} moves through the layers:

```
a = self.linear(x)
p = self.softmax(a)
return p
```

Initialize the model using

```
model = Softmax(n_feat, n_classes)
```

where n_{feat} is the number of data features and n_{classes} is the number of potential class labels.

(c) **Loss Definition:**

Implement the function `cross_entropy(y_pred, y_gt)` which computes and returns the cross entropy between the prediction \mathbf{y}_{pred} and the (one-hot encoded) ground truth class \mathbf{y}_{gt} .

(d) **Optimizer Definition:**

PyTorch provides many different built-in optimization algorithms like Stochastic Gradient Descent (SGD) which can be used to update the weights of your model. Define a SGD-optimizer for the parameters of the initialized model with the learning rate $lr = 0.01$ using:

```
optimizer = torch.optim.SGD(model.parameters(), lr = 0.01)
```

(e) **Train the Model:**

This is done by passing your input data through the model in the forward direction, calculating the loss, and then backpropagating the error by calling `loss.backward()`. Then, you update the weights using the optimization algorithm by calling `optimizer.step()`. The complete training algorithmic is:

```
Loss = []
epochs = 100
len_traindata = #TODO
for epoch in range(epochs):
    for i in range(len_traindata):
        x = torch.from_numpy(X_train[i])
        y = torch.from_numpy(y_train[i])
        optimizer.zero_grad()
        y_pred = model(x)
        loss = cross_entropy(y_pred, y)
        loss.backward()
        optimizer.step()
        Loss.append(loss.detach())
print("Model training is finalized!")
```

Perform the training of the model and plot the Loss curve, where the x-axis represents the training steps and the y-axis the according loss values.

(f) **Evaluating the Model:**

Once you have trained the model, the model performance is evaluated using the test data. Therefore, set the model in evaluation mode using: `model.eval()` before doing so. Similar to the training routine, iterate (only once!) over the test data and compute the average cross entropy loss.